

Evaluating Classifiers




Agenda

- Introduction
- Evaluation Strategies
 - Re-substitution
 - Leave one out
 - Hold out
 - Cross Validation
 - Bootstrap Sampling
- Classifiers Evaluation Metrics
 - Binary classification confusion matrix
 - Sensitivity and specificity
 - ROC curve
 - Precision & Recall
 - Multiclass classification




Introduction

- Evaluation aims at selecting the most appropriate learning schema for a specific problem
- We evaluate its ability to generalize what it has been learned from the training set on the new unseen instances



Evaluation Strategies

- Different strategies to split the dataset into training ,testing and validation sets.
- Error on the training data is not a good indicator of performance on future data, *Why?*
 - Because new data will probably **not** be *exactly* the same as the training data!
 - The classifier might be fitting the training data too precisely (called *Over-fitting*)




Re-substitution

- Testing the model using the dataset used in training it
- Re-substitution error rate indicates only how good /bad are our results on the TRAINING data; expresses some knowledge about the algorithm used.
- The error on the training data is NOT a good indicator of performance on future data since it does not measure any not yet seen data.



Leave one out

- leave one instance out and train the data using all the other instances
- Repeat that for all instances
- compute the mean error



Hold out

- The dataset into two subsets use one of the training and the other for validation (In Large datasets)
- Common practice is to **train** the classifiers using **2 thirds** of the dataset and **test** it using the **unseen third**.
- ***Repeated holdout method:*** process repeated several times with different subsamples
- error rates on the different iterations are averaged to give an overall error rate




K-Fold Cross Validation

- Divide the data to k equal parts (folds),
- One fold is used for testing while the others are used in training,
- The process is repeated for all the folds and the accuracy is averaged.

- 10 folds are most commonly used

- Even better, repeated Cross-validation

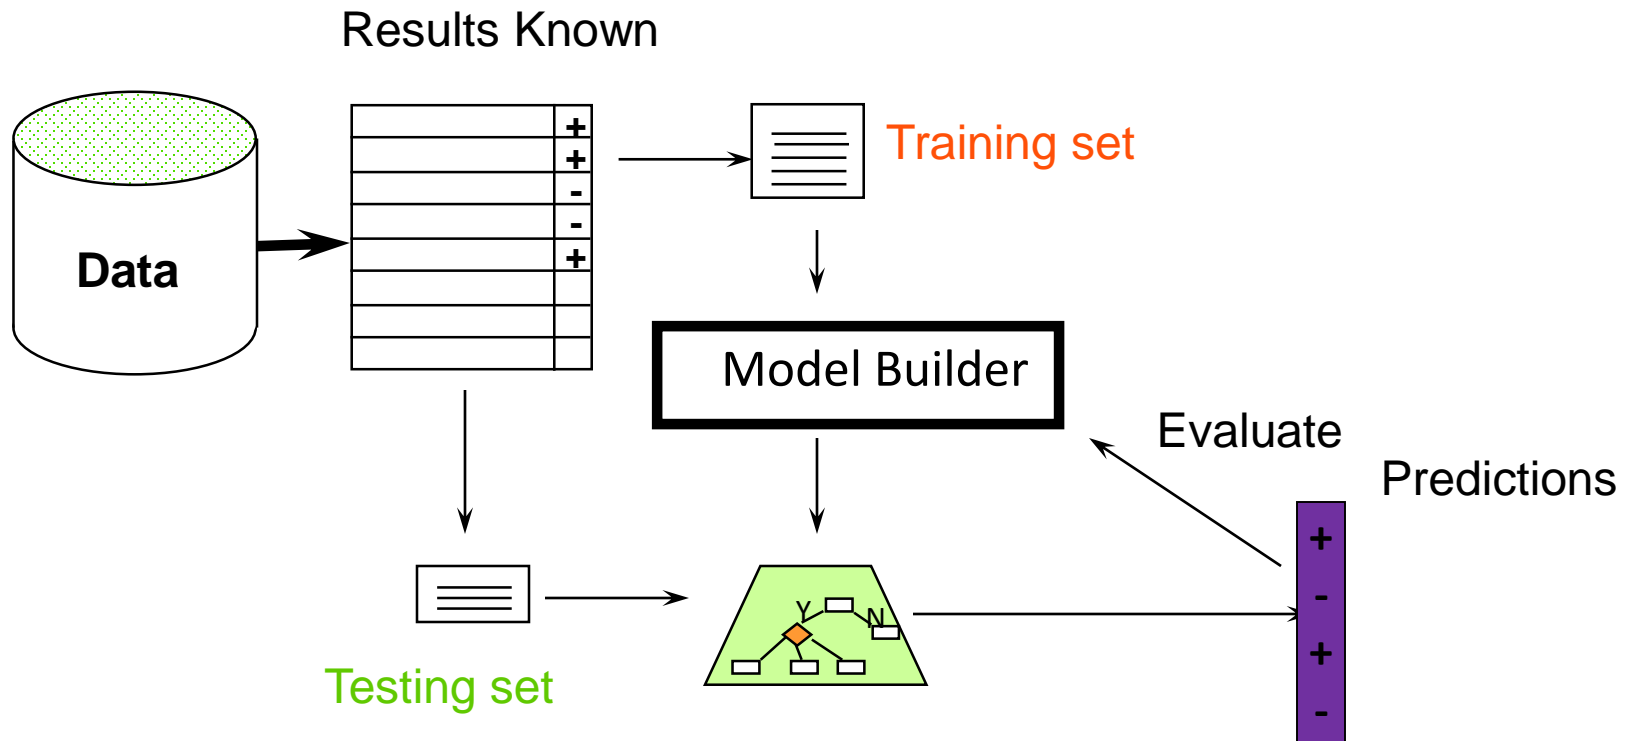


Bootstrap sampling

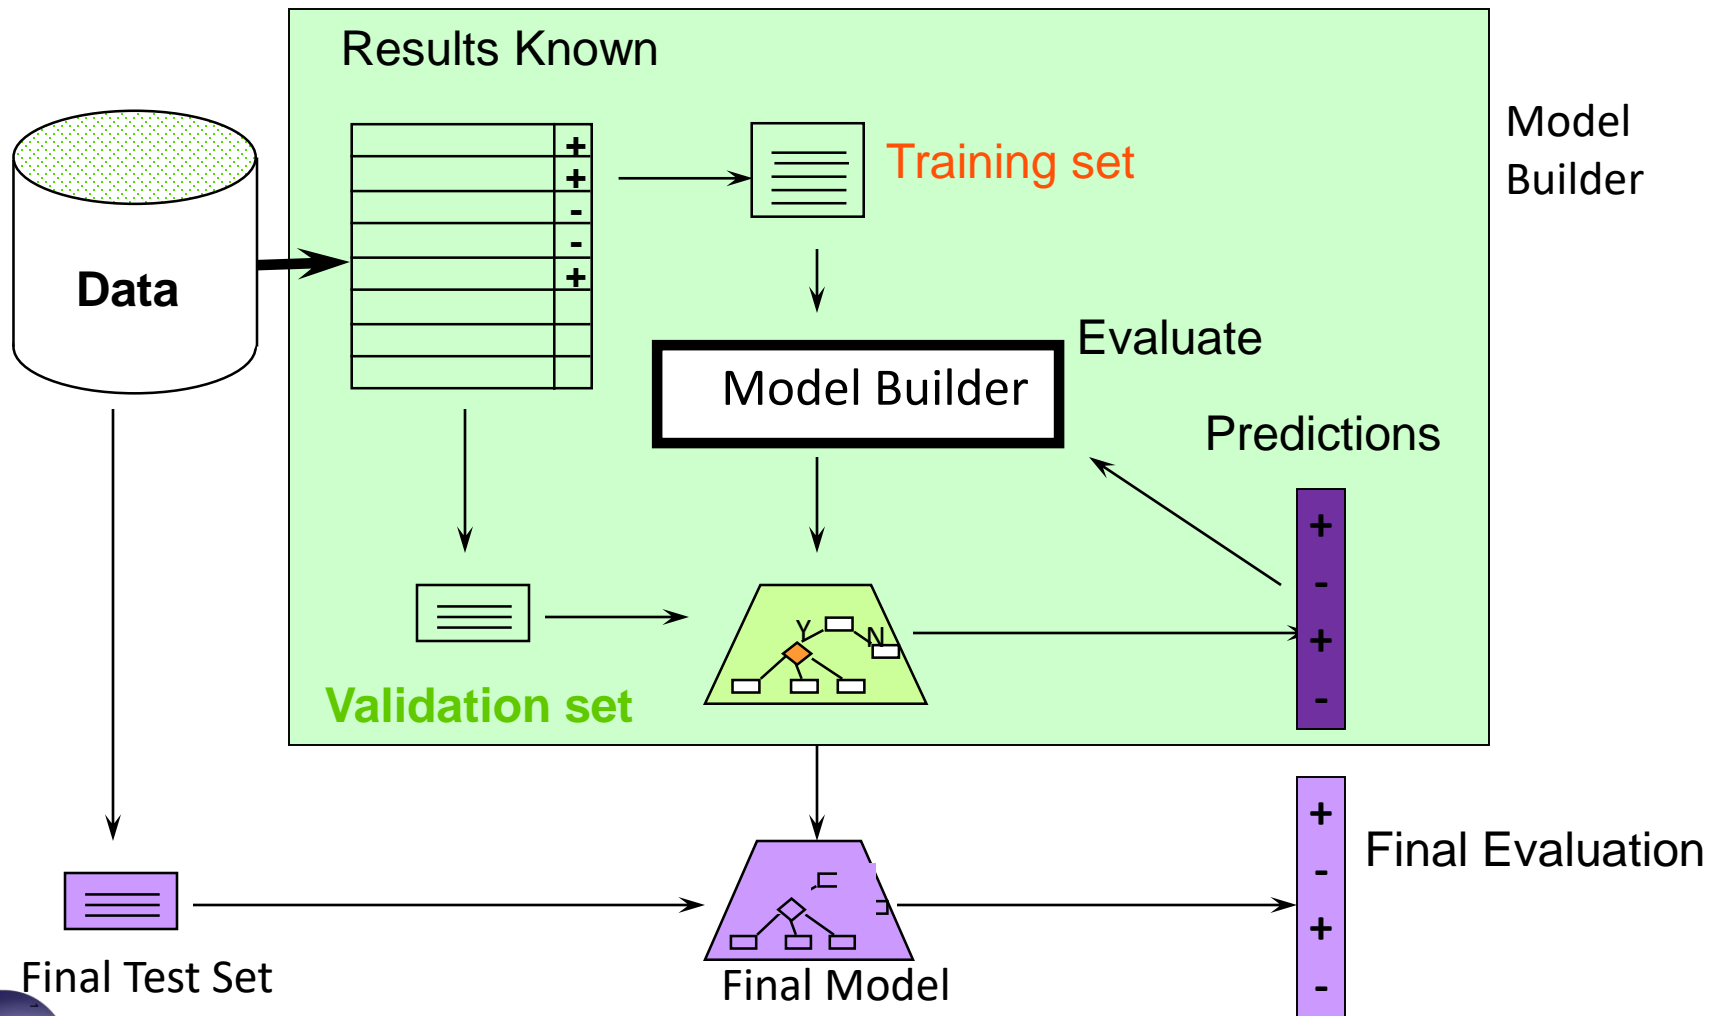
- CV uses sampling without replacement
 - The same instance, once selected, can not be selected again for a particular training/test set
- The bootstrap uses sampling with replacement to form the training set
 - Sample a dataset of n instances n times with replacement to form a new dataset of n instances
 - Use this data as the training set
 - Use the instances from the original dataset that don't occur in the new training set for testing



Are Train/Test sets enough?



Train, Validation, Test split



Metrics

It is extremely important to use **quantitative metrics** for evaluating a machine learning model

- Until now, we relied on the **loss (objective) function value** for regression and classification
- Other metrics can be used to **better evaluate** and understand the model
- **For classification**
 - ✓ Accuracy/Precision/Recall/F1-score, ROC curves,...
- **For regression**
 - ✓ Normalized RMSE, Normalized Mean Absolute Error (NMAE),...



RMSE Metric for Regression

$$RMSE = \sqrt{\sum_{i=1}^n \frac{(\hat{y}_i - y_i)^2}{n}}$$

$\hat{y}_1, \hat{y}_2, \dots, \hat{y}_n$ are predicted values

y_1, y_2, \dots, y_n are observed values

n is the number of observations



Normalized RMSE

$$NRMSE_m(\mathbf{y}, \hat{\mathbf{y}}) = \frac{RMSE(\mathbf{y}, \hat{\mathbf{y}})}{\frac{1}{n} \sum_{i=1}^n |y_i|} = \frac{RMSE(\mathbf{y}, \hat{\mathbf{y}})}{mean(|\mathbf{y}|)}$$



Mean Absolute Error (MAE)

$$MAE(\mathbf{y}, \hat{\mathbf{y}}) = \sum_{i=1}^n \frac{|y_i - \hat{y}_i|}{n}$$



NMAE

$$NMAE(\mathbf{y}, \hat{\mathbf{y}}) = \frac{MAE(\mathbf{y}, \hat{\mathbf{y}})}{\frac{1}{n} \sum_{i=1}^n |y_i|} = \frac{MAE(\mathbf{y}, \hat{\mathbf{y}})}{mean(|\mathbf{y}|)}$$



Evaluation Metrics For Classifiers


- Accuracy:

$$\text{Accuracy} = \frac{\text{number of correctly classified instances}}{\text{total number of instances}} \times 100$$

- It assumes equal cost for all classes
- Misleading in unbalanced datasets
- It doesn't differentiate between different types of errors
- Example -unbalanced classes
 - Cancer Dataset: 10000 instances, 9990 are **normal**, 10 are **ill**, If our model classified all instances as **normal** accuracy will be 99.9 %
 - Medical diagnosis: 95 % healthy, 5% disease.
 - e-Commerce: 99 % do not buy, 1 % buy.
 - Security: 99.999 % of citizens are not terrorists.



- imbalanced (unbalanced) datasets



- Other examples : Spam classification




Binary classification Confusion Matrix

		predicted	
		negative	positive
actual examples	negative	<i>a</i> TN - True Negative correct rejections	<i>b</i> FP - False Positive false alarms type I error
	positive	<i>c</i> FN - False Negative misses, type II error overlooked danger	<i>d</i> TP - True Positive hits



Binary classification Confusion Matrix



Confusion Matrix



Precision and recall

Suppose that $y = 1$ in presence of a **rare class** that we want to detect

Precision (*How much we are precise in the detection*)

Of all patients where we classified $y = 1$, what fraction actually has the disease?

$$\frac{\text{True Positive}}{\# \text{ Estimated Positive}} = \frac{\text{True Positive}}{\text{True Positive} + \text{False Positive}}$$

Recall (*How much we are good at detecting*)

Of all patients that actually have the disease, what fraction did we correctly detect as having the disease?

$$\frac{\text{True Positive}}{\# \text{ Actual Positive}} = \frac{\text{True Positive}}{\text{True Positive} + \text{False Negative}}$$

$$\text{Success rate} = \frac{TP + TN}{TP + TN + FP + FN}$$

Confusion matrix

		Actual class	
		1 (p)	0 (n)
Estiamted class	1 (Y)	True positive (TP)	False positive (FP)
	0 (N)	False negative (FN)	True negative (TN)



F-measure

- The ***F-measure*** is the harmonic-mean (average of rates) of precision and recall and takes account of both measures.

$$F\ measure = \frac{1}{\frac{1}{Recall} + \frac{1}{Precision}} = \frac{2 \times Precision \times Recall}{Precision + Recall} = \frac{2 \times TP}{2 \times TP + FP + FN}$$

- It is biased towards all cases except the true negatives



Example:

Dataset:

- Contains 39 instances, 10 attributes
- The class labels are “negative, positive”
- 22 positive & 17 negative instances.

Classifier used: J48 – 10 folds cross validation

Confusion Matrix:

Classified as →	Positive	Negative
Positive	22	0
Negative	17	0

Classifier Accuracy = $\frac{22}{39} \times 100 = 56.4\%$

- TP= 22
- TN= 0
- FP= 17
- FN= 0
- The area under ROC curve: 0.5 in both cases, cause the TP rate = FP rate.
- Precision & Recall
 - Recall = $\frac{22}{22+0} = 1$
 - Precision = $\frac{22}{22+17} = 0.564$
- The **F-measure** = $\frac{2 \times 22}{2 \times 22 + 17 + 0} = 0.7213$



F1-score

It is usually better to compare models by means of one number only. The **F1 – score** can be used to combine precision and recall

	Precision(P)	Recall (R)	Average	F ₁ Score
Algorithm 1	0.5	0.4	0.45	0.444
Algorithm 2	0.7	0.1	0.4	0.175
Algorithm 3	0.02	1.0	0.51	0.0392

The best is Algorithm 1

Average says not correctly that Algorithm 3 is the best

$$\text{Average} = \frac{P + R}{2}$$

$$F_{\text{score}} = 2 \frac{P \cdot R}{P + R}$$

- $P = 0$ or $R = 0 \Rightarrow F_{\text{score}} = 0$
- $P = 1$ and $R = 1 \Rightarrow F_{\text{score}} = 1$



Precision vs Recall

- To summarize:
 - A system with high Precision might leave out some actual positives, but what it intends to return is high accuracy on the positive class.
 - On the other hand, a system with high Recall might give you numerous misclassifications of actual negatives, but it almost always will correctly classify the actual positives from the dataset.
 - Which one to choose is a tough decision and entirely depends on the problem you are using machine learning for.



Precision vs Recall

- Generally speaking, optimizing Precision typically reduces Recall and vice versa.
- Therefore, many applications consider a metric derived from both Precision and Recall (called the F-score) to measure the performance of a machine learning situation.



Precision vs Recall-Example 1

- Say you are building a course recommender system. The idea is to recommend courses to students based on their profiles.
- A student's time is crucial; thus, you don't want to waste their valuable time recommending courses that they may not like or are irrelevant to them.
- What would you optimize? Precision or Recall?



Precision vs Recall-Example 1

➤ Answer:

- Precision is what you should prefer optimizing for. It is okay not to recommend good courses.
- However, what your system recommends should be very high quality as students should not spend time watching courses that are irrelevant to them.



Precision vs Recall-Example 2

- Next, assume that you want to shortlist candidates for a job opening and see if they should be selected for an interview.
- As your organization is looking for talented candidates, you don't want to miss out on a candidate that can be a potential hire.
- What would you optimize for now? Precision or Recall?



Precision vs Recall-Example 2


➤ Answer:


- Having inexperienced or irrelevant candidates won't hurt as long as you catch all the relevant and talented ones.
- Therefore, Recall is the metric you should optimize for. If you instead optimize for Precision, there is a possibility of excluding talented candidates, which is not desired.




ROC Curves

- Stands for “receiver operating characteristic”
- Used in signal detection to show tradeoff between hit rate and false alarm rate over noisy channel





TP	FP
FN	TN
1	1




Identify a threshold in your classifier that you can shift.


Plot ROC curve while you shift that parameter.




Stochastic model, Gaussians, equal variances




ROC curve




Stochastic model, Gaussians, equal variances




ROC curve



Stochastic model, Gaussians, equal variances




ROC curve



ROC

ROC curves are a very general way to **represent and compare** the performance of different models (on a binary classification task)



Observations

- (0,0) : classify always negative
- (1,1) : classify always positive
- Diagonal line: random classifier
- Below diagonal line: worse than random classifier
- Different classifiers can be compared



ROC curve (Cnt'd)

- Accuracy is measured by *the area under the ROC curve*. (AUC)
An area of 1 represents a perfect test; an area of .5 represents a worthless test:
 - .90-1 = excellent
 - .80-.90 = good
 - .70-.80 = fair
 - .60-.70 = poor
 - .50-.60 = fail



Multiclass classification

- For **Multiclass prediction** task, the result is usually displayed in confusion matrix where there is a row and a column for each class,
 - Each matrix element shows the number of test instances for which the actual class is the row and the predicted class is the column
 - Good results correspond to large numbers down the diagonal and small values (ideally zero) in the rest of the matrix

Classified as →	a	b	c
A	TP_{aa}	FN_{ab}	FN_{ac}
B	FP_{ab}	TN_{bb}	FN_{bc}
C	FP_{ac}	FN_{cb}	TN_{cc}



Multiclass classification (Cont'd)

- For example in three classes task $\{a, b, c\}$ with the confusion matrix below, if we selected a to be the class of interest then

True positives for class $a = TP_{aa}$

True Negatives for class $a = TN_{cc} + TN_{bb}$


False Positives for class $a = FP_{ab} + FP_{ac}$

False Negatives for class $a = FN_{ab} + FN_{ac}$


- Note that we don't care about the values $(FN_{cb} \text{ \& } FN_{bc})$ as we are considered with evaluating how the classifier is performing with class a , so the misclassifications between the other classes is out of our interest.




Confusion matrix For Multi-class Classification



Confusion matrix For Multi-class Classification



Confusion matrix For Multi-class Classification



Confusion Matrix for multiclass Classifications



Multiclass classification (Cont'd)

- To calculate overall model performance, we take their weighted average to evaluate the overall performance of the classifier.
- Averaged per category (macro average) :
 - Gives equal weight to each class, including rare ones

$$Recall_{macro} = \frac{\sum_{i=1}^M Recall_i}{M}, \text{ where } M \text{ is number of classes}$$

$$Precision_{macro} = \frac{\sum_{i=1}^M Precision_i}{M}, \text{ where } M \text{ is number of classes}$$



Multiclass classification (Cont'd)

➤ Micro Average:

- Obtained from true positives (TP), false positives (FP), and false negatives (FN) for each class, and F-measure is the harmonic mean of micro-averaged precision and recall
- Micro average gives equal weight to each sample regardless of its class.
- They are dominated by those classes with the large number of samples.

$$Recall_{micro} = \frac{\sum_{i=1}^M TP_i}{\sum_{i=1}^M TP_i + \sum_{i=1}^M FN_i}$$

$$Precision_{micro} = \frac{\sum_{i=1}^M TP_i}{\sum_{i=1}^M TP_i + \sum_{i=1}^M FP_i}$$



Example: Area under ROC (AUC)

Dataset	J48	BayesNaive	Multi layer Perceptron	SVM	K Star	CART
approach_1	0.57	0.41	0.52	0.45	0.66	0.58
approach_2	0.48	0.43	0.66	0.5	0.64	0.58
approach_3_10+threshold	0.61	0.63	0.44	0.48	0.53	0.48
approach_3_15+threshold	0.55	0.73	0.59	0.59	0.65	0.54
approach_3_7+threshold	0.62	0.47	0.32	0.39	0.39	0.54
approach_4_10_threshold	0.63	0.49	0.47	0.44	0.55	0.53
approach_4_15_threshold	<u>0.76</u>	0.43	0.44	0.53	0.65	0.53
approach_4_7_threshold	0.35	0.46	0.48	0.51	0.44	0.48



Example: F-Measure

Dataset	J48	BayesNaive	Multi layer Perceptron	SVM	K Star	CART
approach_1	0.4	0.21	0.4	0.03	0.51	0.4
approach_2	0.06	0.31	0.45	0	0.54	0.44
approach_3_10+threshold	0.59	0.52	0.46	0.42	0.49	0.5
approach_3_15+threshold	0.56	0.61	0.54	0.52	0.48	0.53
approach_3_7+threshold	0.61	0.46	0.32	0.37	0.38	0.62
approach_4_10_threshold	0.5	0.36	0.32	0.23	0.29	0.18
approach_4_15_threshold	0.67	0.24	0.35	0.38	0.33	0.18
approach_4_7_threshold	0.07	0.26	0.31	0.27	0.27	0.04

